

# Inf2C - Computer Systems

## Lecture 17-18

### I/O

---

Boris Grot

School of Informatics  
University of Edinburgh



# Previous lecture: Exceptions & processor mgmt

---

- Exceptions: interrupt normal program flow and require servicing by the CPU
  - Internal “traps” (e.g., syscall)
  - external “interrupts” (e.g., keyboard click)
- Exception handling
- Processor protection modes (user vs kernel)
- Processor management via time-sharing

# Examples of I/O Devices

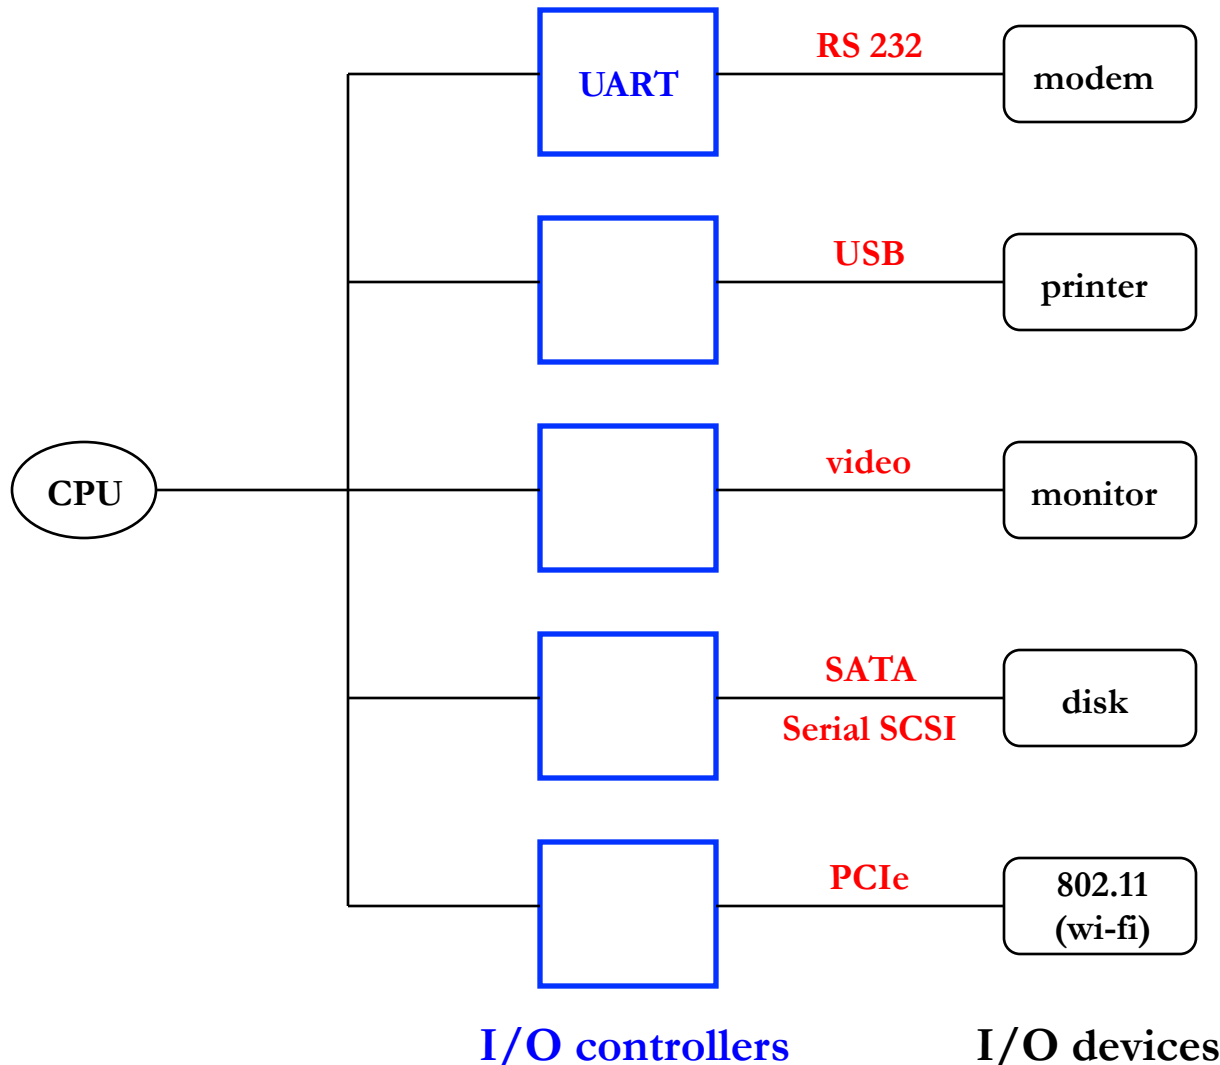
---

Device	Behaviour	Partner	Data Rate (Mbit/sec)
Keyboard	Input	Human	0.001
Mouse	Input	Human	0.004
Voice input	Input	Human	0.26
Laser printer	Output	Human	3.2
Graphics	Output	Human	800-8000
Magnetic disk	Storage	Machine	800-3000
Network/LAN	Input or output	Machine	100-40,000 (40Gbit/sec)

**Retina display:  
>16Gbit/s**

**Wi-fi: 10-300 Mbit/s**

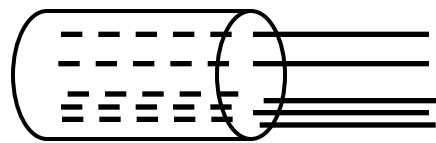
# I/O Controllers & Devices



# Example: RS232 Serial Interface

---

- I/O controller: UART
  - Stands for Universal Asynchronous Receiver Transmitter
- Used for modems and other serial devices
- Physical Implementation:
  - 2 signal wires (one for each direction) + ground reference + status signals



**Tx data wire (CPU to I/O device)**

**Rx data wire (I/O device to CPU)**

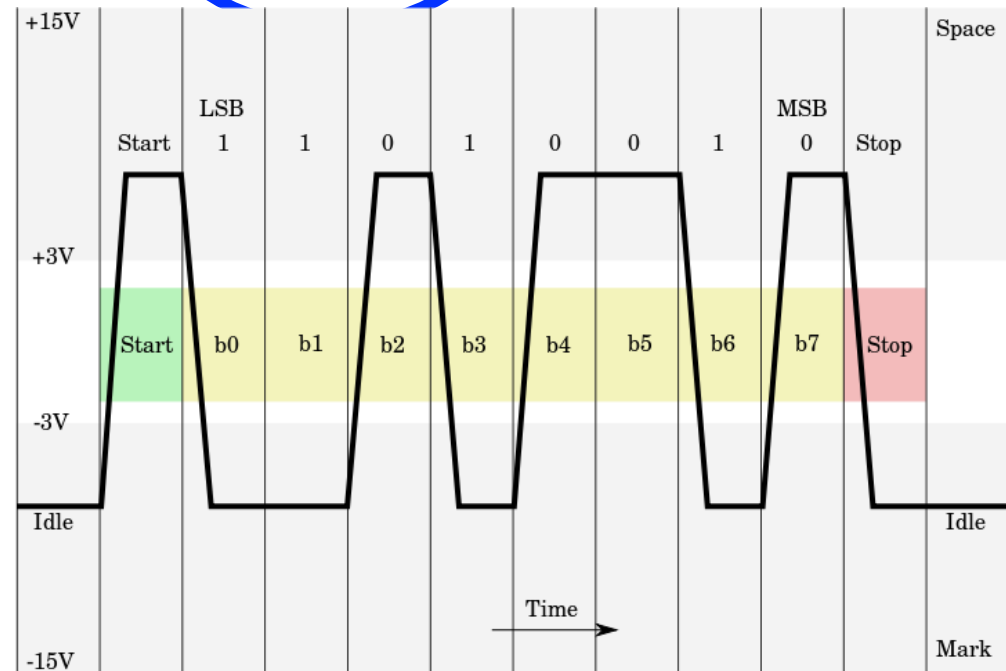
**GROUND wire, status signals**

# RS232 modem: bits and wires




*Oscilloscope trace of voltage levels for an ASCII "K" character (0x4B) with 1 start bit, 8 data bits, 1 stop bit.*

*Source: wikipedia*



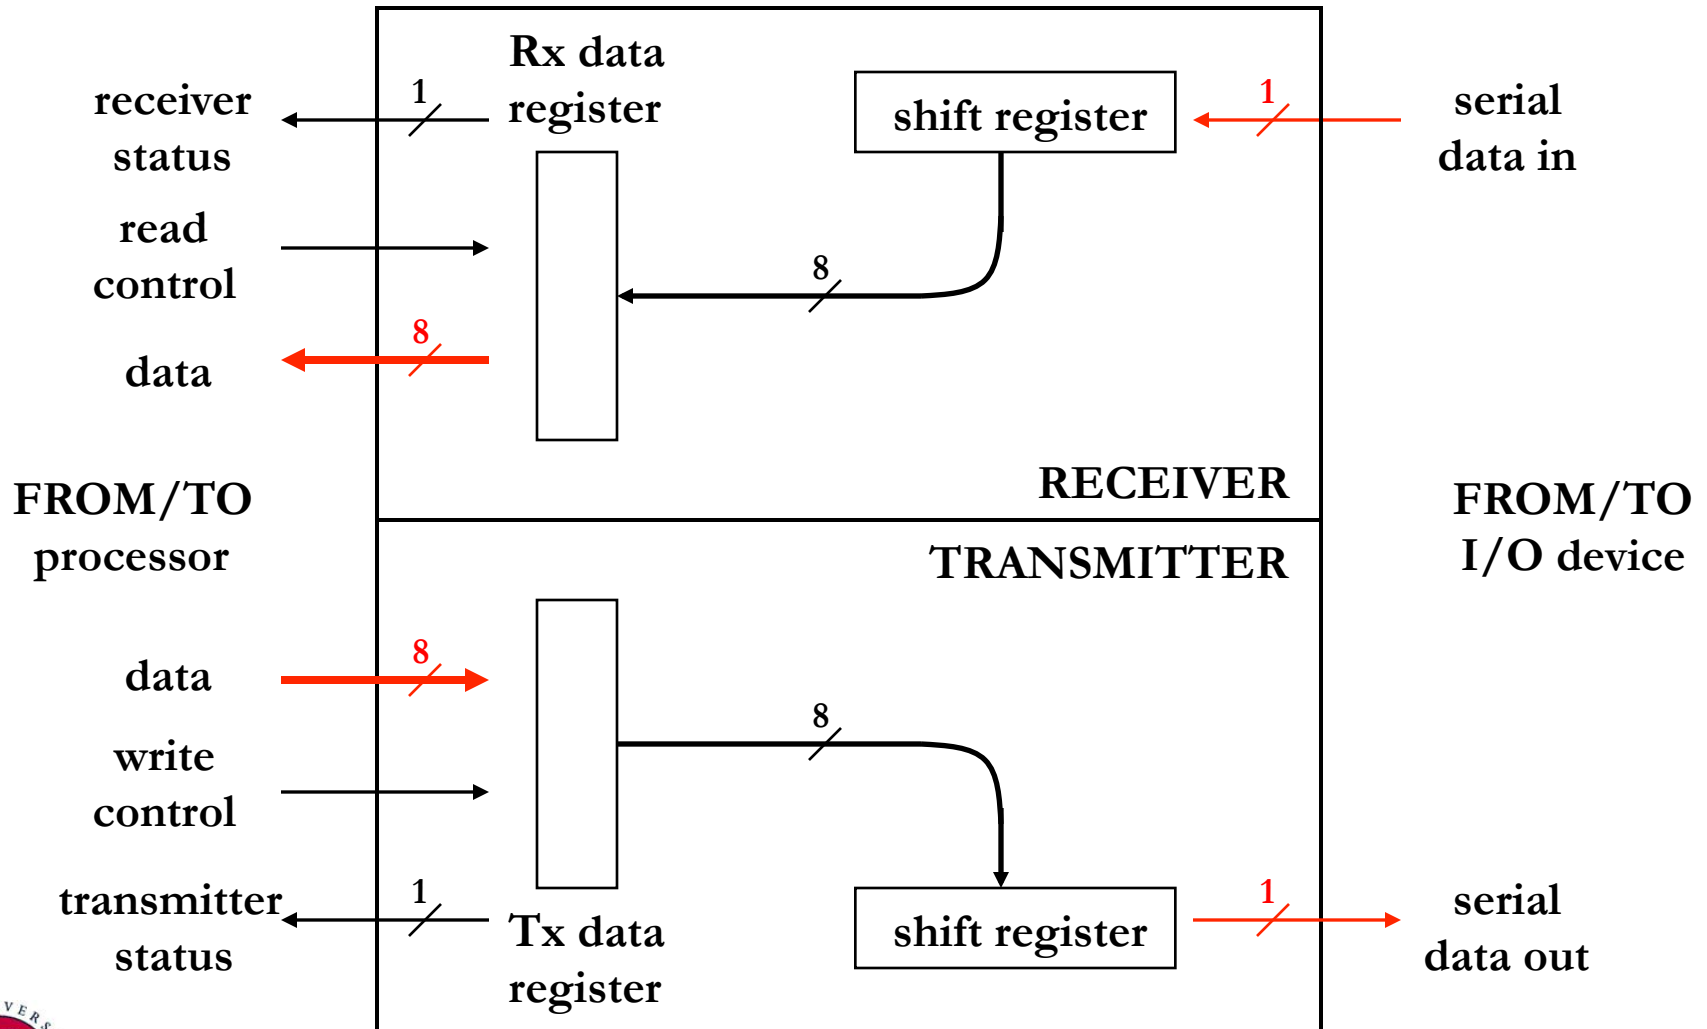
# Example: RS232 Serial Interface

- Encoding:
  - 1 character = 10 or 11 bits (including signaling)
  - Idle state is represented by a constant “1”



- Parity: for detection of transmission errors
  - odd → total number of 1's (including parity bit) is odd
  - even → total number of 1's is even

# UART Controller





# Connecting CPU and I/O Controllers

---

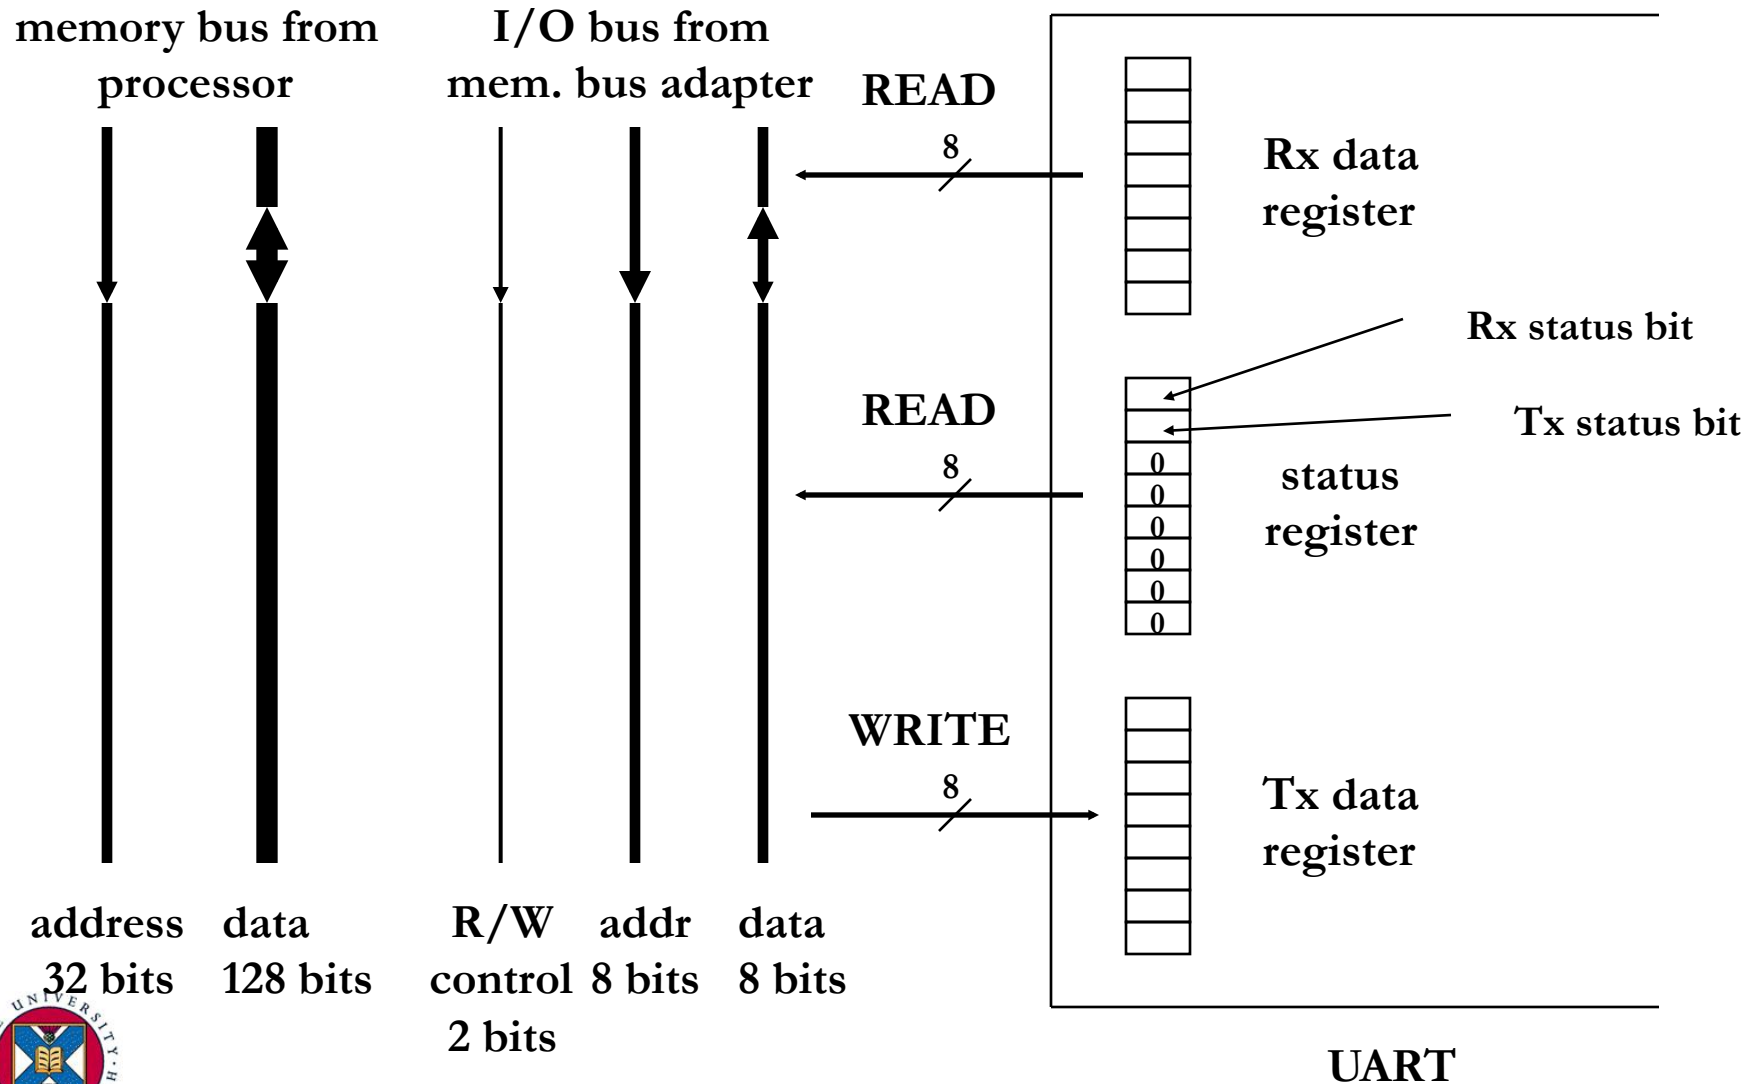
- Option 1: connect the I/O Tx and Rx registers directly into some special CPU I/O registers → not flexible
- Option 2: keep I/O registers in separate I/O controller and connect CPU to I/O controller through special I/O bus → expensive, not flexible

I/O bus:

- data lines (8 bits)
- control lines (READ and WRITE signals),
- address lines (some few bits) → each I/O controller is assigned a range of addresses for its registers

Data is accessed through special I/O loads and stores

# I/O via a Dedicated I/O Bus

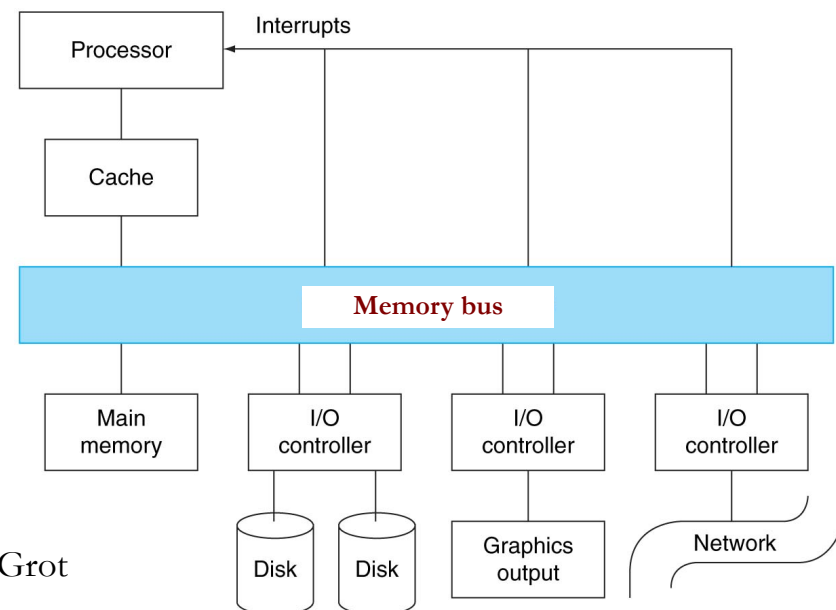


# Connecting CPU and I/O Controllers

- Option 3: keep I/O registers in I/O controller and connect the CPU to I/O controller through memory bus

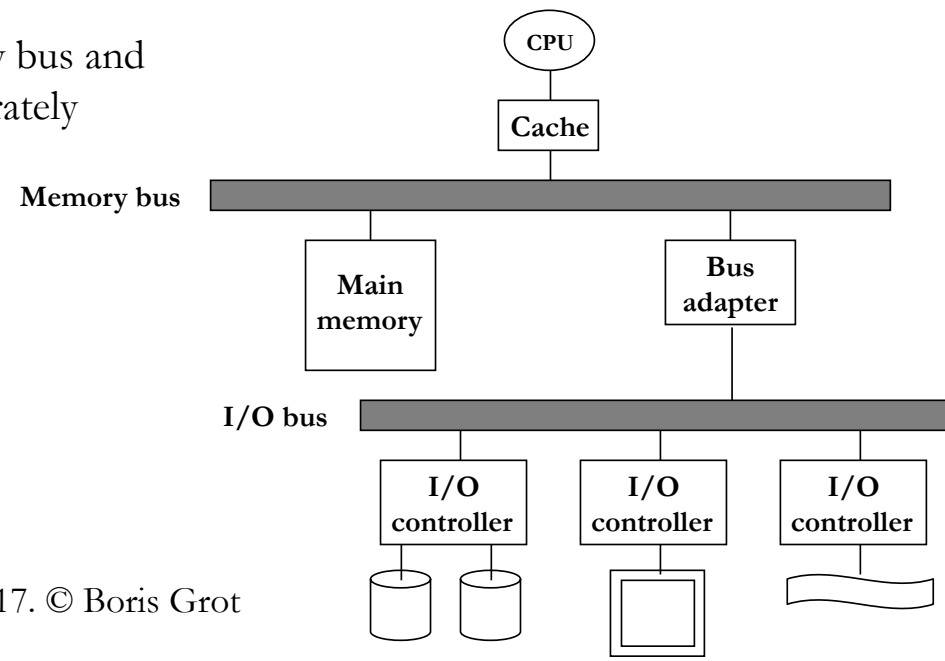
## Memory mapped I/O:

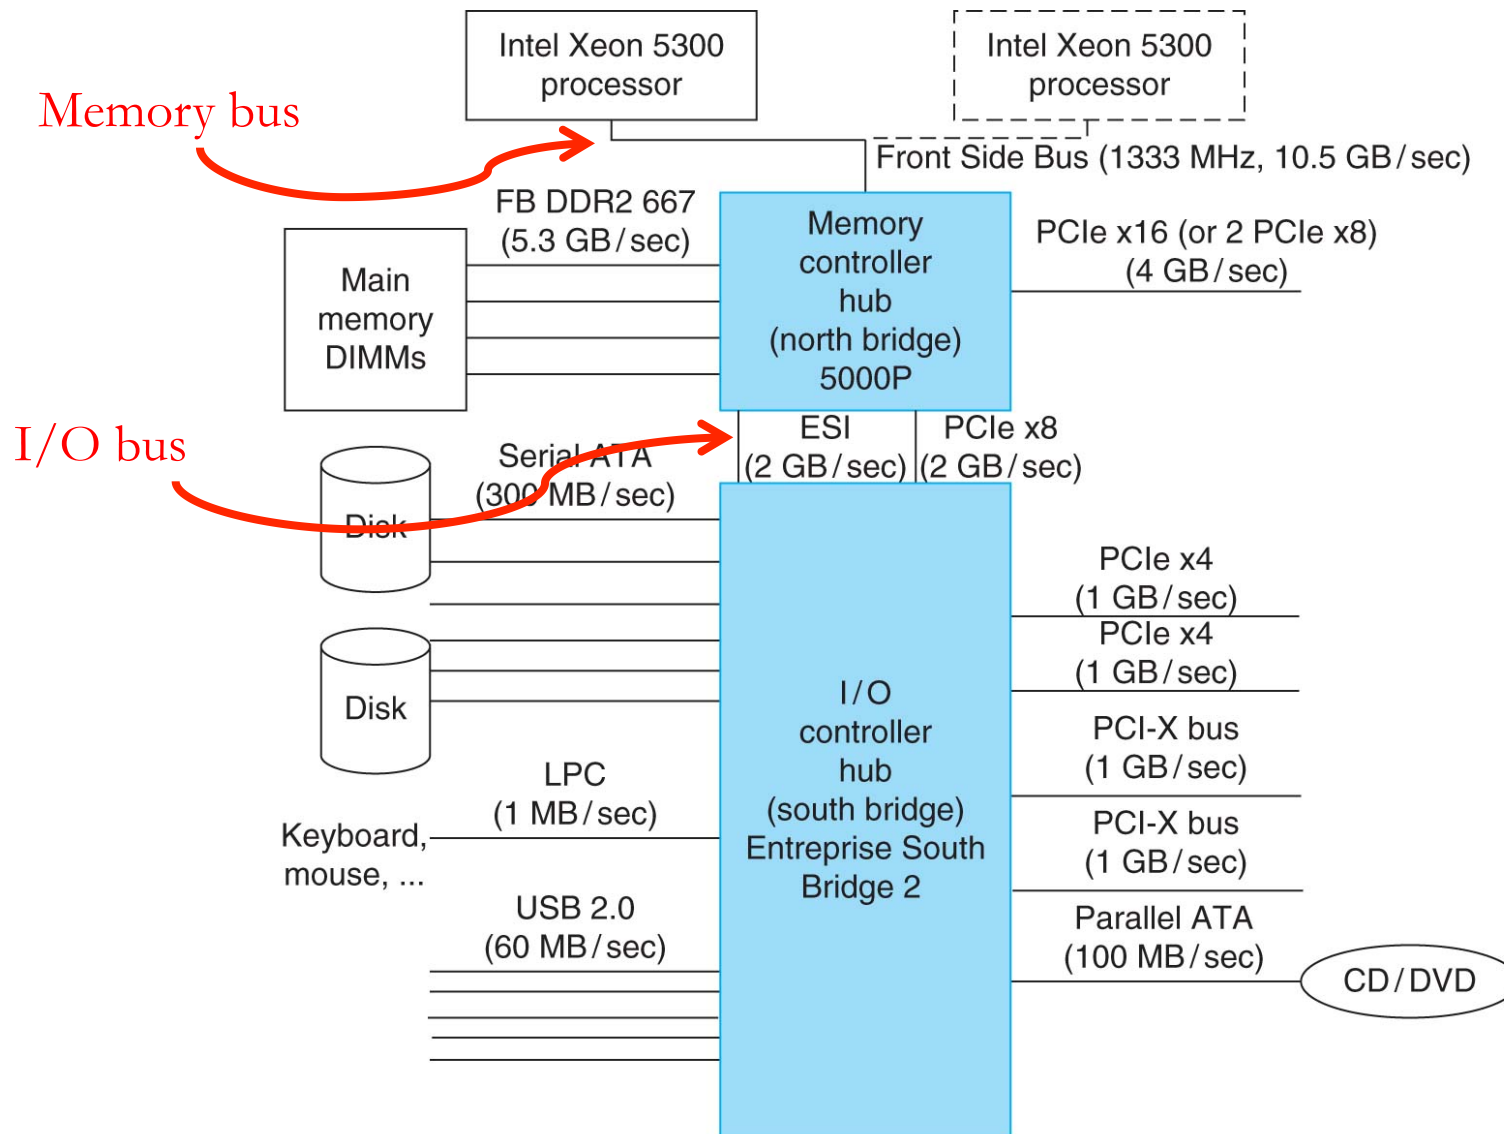
- I/O controller registers (data and control) are mapped to a dedicated portion of memory
- Good news: **accessed with regular load and store instructions**
- Bad news: Takes bus bandwidth away from CPU-memory



# Connecting CPU and I/O Controllers

- Option 4: connect I/O controllers to I/O bus and the I/O bus to the memory bus through a bus adapter
  - Off-load the I/O from the memory bus: multiple I/O devices appear as a single device to the memory bus)
  - Pros:
    - better performance: slow & narrow I/O bus doesn't clog up the fast memory bus
    - higher flexibility: add/remove devices without impacting the high-performance processor-memory interface
    - Modularity and extensibility: memory bus and I/O bus technology can evolve separately





Organization of the Memory & I/O system on an Intel server using the Intel 5000P chip set.

# Polling and interrupt-based I/O

---

How do we check I/O status (e.g., key clicked)?

- Option 1: Polling
  - User process calls OS at regular intervals to check status of I/O operation
  - Wasteful for events that happen very infrequently (e.g., keyboard clicks)
- Option 2: Interrupt
  - I/O controller interrupts user process to signal an I/O event
  - Heavy-weight (break out of user code into kernel mode)



# USB example to bring things together

---

- USB devices don't generate interrupts
  - Keeps cost low
- Computer can't afford to keep polling the multitude of USB ports
  - Would lose too much performance for nothing
  - Not polling often enough not an option: will lose data
- USB controller on the CPU side does the polling (via a simple FSM) and generates an interrupt to inform the CPU as needed
  - This functionality is in the Southbridge on the earlier slide



# Hard Disks



- 1-4 platters per drive, 2 surfaces per platter, 10-50k tracks per surface, 100-500 sectors per track, 512B – 4KB per sector
- Spinning speed:
  - 5400-15000 rpm (90-250 revs per sec)





# Disk Performance

---

- Total time of a disk operation is divided in two parts:
  1. **Access time:** time to get head into position to read/write data  
 $\text{access time} = \text{seek time} + \text{rotational latency}$ 
    - **Seek time:** time to move head to appropriate track ( $< 10\text{ms}$ )
    - **Rotational latency:** time to wait for appropriate sector to arrive underneath the head ( $< 10\text{ms}$ )  
Dependent on spinning speed
  2. **Transfer time:** time to move data to/from disk  
 $\text{transfer time} = \text{time to transfer 1 byte} * \text{number of bytes of data}$ 
    - Dependent on both spinning speed and recording density
    - 75-125 MB/s (changing very slowly – limited by mechanics)

# Disk Controllers

---


- Disk controller inside disk unit → responsible for all mechanical operation of disk + interface with CPU
- I/O registers:
  - Exchange data and status words between CPU and disk controller
  - Command register → tells disk controller what to do next

e.g. Seek  $n$

Read Sector  $m$

Write Sector  $m$

Format Track



**High-level commands**

# Using a Disk Controller

---

- Step 1: user program requests data from a file
- Step 2: OS file system determines sector(s) to be accessed
- Step 3: OS disk handler issues **Seek** command and CPU goes to work on some other process (multi-tasking)
- Step 4: I/O controller interrupts CPU to signal completion of seek
- Step 5: OS disk handler issues **Read Sector** command and CPU goes to work on some other process
- Step 6: I/O controller interrupts CPU to signal data ready
- Step 7: OS disk handler transfers sector data from disk controller to memory
  - This is a slow-running *for* loop that transfers data word-by-word via the I/O and memory buses

Step 8: go to step 3 or 5 and repeat until all data transferred

# Problem with the Interrupt Approach

---

- CPU (through OS) has to issue individual commands to read every sector from disk
- Frequent interrupts detrimental to processor performance
  - Switch to OS, then switch back to the application (Lecture 13)
- CPU moving the data from the disk controller via the slow I/O bus results in highly inefficient CPU utilization
- Solution: Direct Memory Access (DMA)

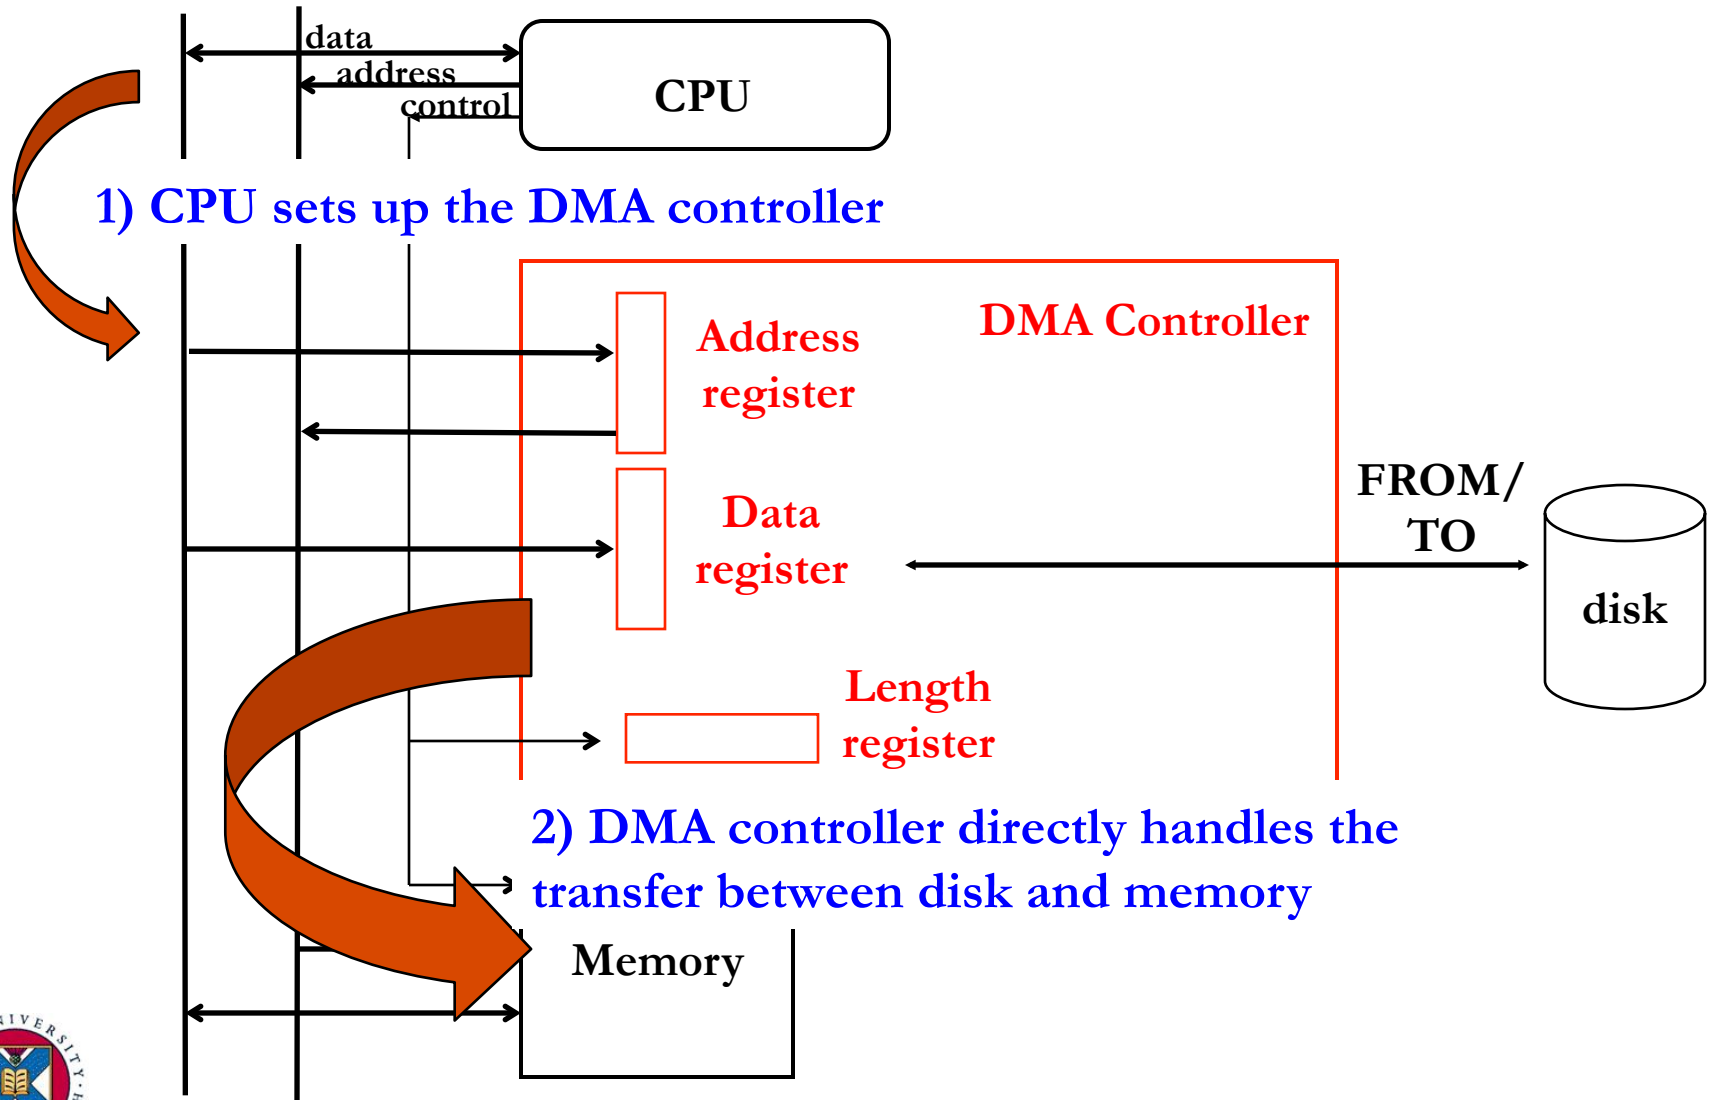
# Direct Memory Access (DMA)

---

- DMA controller: stateful device that sits on the memory bus and can independently transfer data between memory and disk
  - Setup by the processor for each transfer using memory-mapped registers inside the DMA controller
- DMA registers:
  - Address register → position in memory of next data to be read/written
  - Data register → temporary storage for data to be transferred
  - Length register → number of bytes remaining to be transferred
- Technical write-up of how this works for x86:  
<http://www.brokenthorn.com/Resources/OSDev21.html>



# DMA Organization



# DMA Operation

---

- Step 1, 2: user program requests data, OS determines location of data on disk
- Step 3: OS disk handler issues **Seek** command and sets up DMA registers (address, length); CPU goes to work on another process
- Step 4, 5: I/O controller interrupts CPU, OS disk handler issues **Read Sector** command
- Step 6: I/O controller informs DMA controller that data is ready (no need to interrupt CPU)
- Step 7: DMA controller transfers data into memory; length register is decremented until all data is moved (advanced DMA controllers can access multiple tracks with a single operation)
- Step 8: DMA controller interrupts CPU to inform completion of DMA operation



# Bus Arbitration

---

- DMA and CPU connect to memory bus → access must be somehow arbitrated to avoid conflicts
- Solution: additional logic (**bus arbiter**)
  - Authorizes CPU or DMA controller to access memory at any given time
- Two new wires on the memory bus:
  - Bus Request → asserted by the DMA controller when it requires the bus
  - Bus Grant → asserted by the CPU when it is not using the bus and thus the DMA controller can use it
  - In case of conflicting requests in the same cycle, CPU usually has priority



*That's all Folks!*

