

Inf2C Computer Systems

Tutorial 4, Week 9

Solutions

Boris Grot, Paul Jackson, Stratis Viglas

1. Multi-cycle processor

Discuss the steps in executing the `jal`, the `lw` and the `add` instruction in the multi-cycle datapath presented in Figure 1.

Now assume that an access to the register file takes 6.5ns, an access to the memory takes 100ns and the ALU delay is 6ns. Assume also that the instruction's execution stalls when it waits for a resource to produce a result. If we were to optimize this simple processor, which should be the component we should spend our main design efforts on, what could we do to improve it?

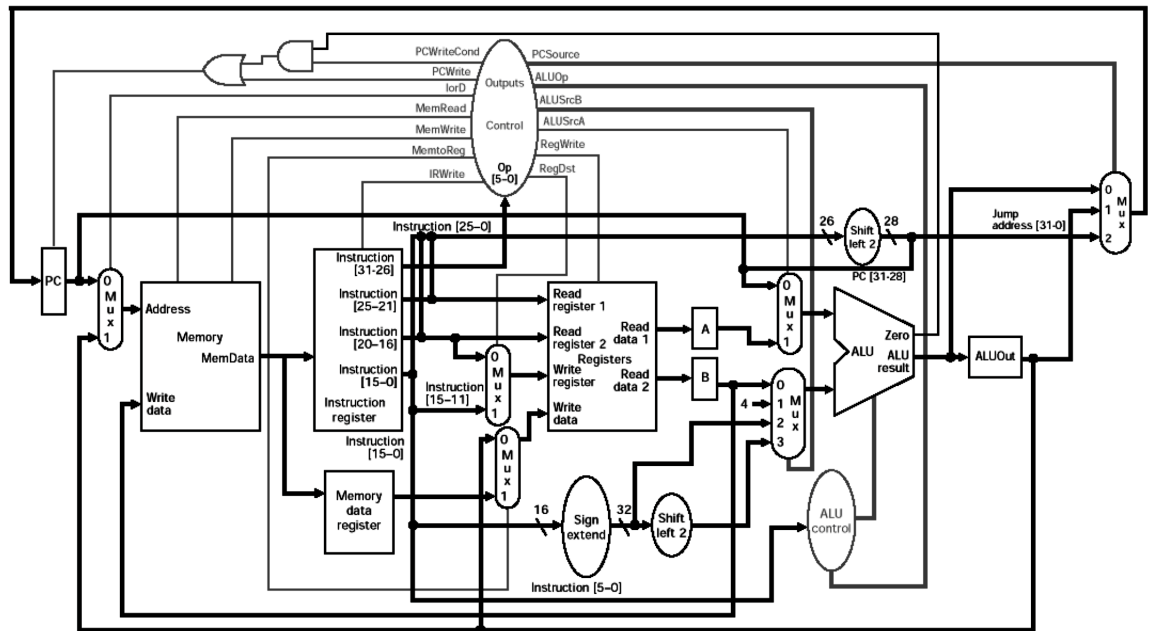


Figure 1: MIPS Datapath

Answer: The first two cycles are the same for all the instructions. During the first cycle, we fetch in the Instruction Register from the memory location indicated by the current PC, and we increment the PC by 4.

In the second cycle, we read the values in the register file registers addressed by bits 26-21 and bits 20-16 of the loaded instruction into registers A and B respectively. We also use the ALU to speculatively compute a branch destination address, though that is not needed by any of the instructions under consideration.

JAL:

JAL (jump and link) requires that we save the current program counter ($PC+4$, where PC is the address of the JAL instruction) to register 31. We can do so in the 3rd cycle with a couple of small modifications: we need the write register multiplexer extended to take 31 as an extra input and the write data multiplexer extended to take the current value in the program counter. At the same time we compute the jump location using the shift left by 2 of the lower 26 bits of the Instruction Register concatenated with the upper 4 bits of the current program counter. We then load this new value into the program counter at the end of the 3rd cycle.

LW:

The LW is a memory instruction, so in the third cycle we have to form the effective address. We use the ALU to add the A register contents to the sign extended Immediate value from the Instruction Register.

In the fourth cycle we load the Memory Data Register MDR with the memory location computed in the previous cycle.

In this final cycle, we load the register indicated by the bits 16 to 20 of the Instruction Register, with the value we have in the MDR.

ADD:

The ADD is a register-register operation. So in the third cycle, we compute the addition using the ALU of the contents of register A and register B.

In the fourth cycle we load the register indicated by the bits 11 to 15 of the instruction register with the computed value.

In the final section, note that the memory operation costs a lot more than the access to the register file and the ALU. Therefore one should try to improve that. The obvious way to do it is to use a cache to provide a fast access for the vast majority of the memory operations.

2. Direct-mapped cache

Consider a simple direct-mapped cache with 4 lines, each containing 2 bytes of data and a tag field of appropriate size for 6-bit addresses. Assume that the cache is initially empty and that memory is byte-addressable.

For the following pattern of cache accesses:

Access Number	Address
1	1 0 1 1 0 0
2	0 0 1 0 1 1
3	0 1 1 0 0 1
4	0 1 0 0 0 1
5	1 0 1 1 0 0
6	0 1 1 0 0 1
7	1 0 0 1 1 0

draw up a table of form

Access Number	Tag	Index	Line 0	Line 1	Line 2	Line 3	Hit/Miss
1							
2							
3							
4							
5							
6							
7							

that summarises the contents of the cache *after* each access. In the *Tag* and *Index* columns, write the decimal values of the tag and index components of each access address. Then, in each row, if Line i is occupied with the memory block with tag t , write the decimal value of t in the column for Line i , and record in the Hit/Miss column an H or an M to indicate whether the access is a hit or a miss. If a line is empty after some access, leave the entry for that line and access number blank.

Answer: Each 6-bit address splits into a 3-bit tag, 2-bit index and 1-bit offset. The desired table is

Access Number	Tag	Index	Line 0	Line 1	Line 2	Line 3	Hit/Miss
1	5	2			5		M
2	1	1		1	5		M
3	3	0	3	1	5		M
4	2	0	2	1	5		M
5	5	2	2	1	5		H
6	3	0	3	1	5		M
7	4	3	3	1	5	4	M

Note how line replacement is insensitive to how recently a line has been accessed. E.g. access 4 removes the line accessed immediately before it.

3. Fully-associative cache

Consider again a cache with 4 lines, each containing 2 bytes of data, and assume addresses are 6 bits wide, but this time assume the cache is fully-associative. With the same pattern of accesses as before, draw up a table of form

Access Number	Tag	Line 0	Line 1	Line 2	Line 3	Hit/Miss
1						
2						
3						
4						
5						
6						
7						

that summarises the contents of the cache *after* each access.

Assume that empty cache lines are filled in order and assume an LRU (least recently used) replacement policy. How would the table change if a FIFO (first in first out) replacement policy were used instead?

Answer: Each 6-bit address splits into a 5-bit tag and 1-bit offset. Desired table is:

Access Number	Tag	Line 0	Line 1	Line 2	Line 3	Hit/Miss
1	22	22				M
2	5	22	5			M
3	12	22	5	12		M
4	8	22	5	12	8	M
5	22	22	5	12	8	H
6	12	22	5	12	8	H
7	19	22	19	12	8	M

If a FIFO replacement policy is used instead, only the row for access 7 changes. It becomes:

Access Number	Tag	Line 0	Line 1	Line 2	Line 3	Hit/Miss
7	19	19	5	12	8	M